

Human vs. Agent Writing

Max Simchowitz

A guide to recognizing and revising agent-like prose. The patterns support revision; they do not establish authorship.

Date: July 25, 2026

Abstract

Agent prose can be grammatically clean and still feel wrong. The usual problem is a missing reader. The answer works in general, but no one seems to need this exact sentence. Ordinary facts get inflated, uncertainty disappears, and every point receives the same weight. This guide explains those habits and shows revisions that preserve the writer's facts and voice. The appendix keeps the complete pattern catalog for reference. A later chapter considers several possible causes, including next-token prediction, generic audience models, mixed training data, code and documentation, and post-training incentives. These are hypotheses, not findings. The patterns are useful for revision; they do not establish authorship.

Contents

1	Patterns are useful for revision	3
2	Meaning and voice come first	3
3	Replace performance with information	4
4	The sentence needs a global focus	5
5	Structure should show priority	6
6	Keep the writer in the prose	8
7	Possible causes of agent-like prose	9
8	A five-pass revision	10
A	Pattern map	12
B	Revision checklist	14
C	Sources and acknowledgments	15

1 Patterns are useful for revision

People write into situations. They anticipate an objection, select the fact that is new to this audience, and keep the uncertainty that cannot honestly be removed. An agent can imitate these choices. Without a clear audience, it often produces a generic answer.

That is the distinction this guide cares about. Grammar, formal vocabulary, and punctuation are weak evidence about the source of a passage. Human writers use formulas; agents can write restrained, specific prose. A clean sentence proves nothing about authorship.

1.1 Name the reader

Start an edit by naming the reader and the consequence of the document. A reviewer checking evidence needs different details from a student learning a mechanism. A collaborator choosing between two methods needs the comparison early. Once that purpose is clear, the editor can judge each paragraph by the new understanding or decision it gives that reader.

This test catches vague management language. “Give each section one job” sounds firm but explains little. “Ensure each section helps the reviewer answer one question” gives the rule a reason.

1.2 Describe the text, not the author

Call a passage *formulaic*, *overstated*, or *poorly sourced*. Those claims can be checked. A claim about the person or system that wrote it needs independent evidence.

Style-based detection produces familiar false positives. Professional editing creates consistent grammar, word processors create curly quotation marks, and fields develop shared vocabularies. A journalist may use many dashes. Look for a cluster of choices that causes a reading problem rather than a lone tell.

2 Meaning and voice come first

A successful edit changes the prose without changing the writer’s claim. Style is subordinate to factual fidelity. Every name, number, date, quotation, citation, and technical distinction in the revision must come from the source or from the writer. When the draft lacks a needed fact, either ask for it or write the honest, less specific sentence.

Fiction is the exception because invention belongs to the task. In factual, technical, legal, and reference writing, a vivid unsupported detail remains a fabrication.

2.1 Uncertainty is part of the claim

Preserve the difference between “demonstrates,” “suggests,” “is consistent with,” and “does not rule out.” A shorter sentence is worse when it turns a conditional result into a universal one. The same principle applies to opinions. A writer’s confidence, hesitation, and mixed response should survive the edit.

Missing evidence calls for a boundary, not a plausible substitute. If the available sources do not establish a founding date, say so when the absence matters or omit the sentence. Do not infer a private life, a motivation, or a causal history from silence.

2.2 Voice survives in the choices

Voice appears in selection as much as style. A writer may keep a technical qualification that interrupts the rhythm because it is true. Another may use a dated joke or pause to correct an earlier thought. Do not regularize those choices merely to make every paragraph sound alike.

When the writer supplies an earlier sample, use it as the strongest evidence for vocabulary, punctuation, and rhythm. A generic style rule should not override a consistent choice in that sample unless the choice harms meaning or violates the publication's requirements.

Voice is often uneven. A short sentence may follow a long one, an aside may reveal the writer's stake, and a passage may remain ambivalent. Keep that unevenness when it belongs to the writer.

2.3 A light edit is often enough

Read the full draft before touching a sentence. Identify the central point and the intended reader. Notice a few voice signals worth preserving. Then fix the actual problems and leave strong sentences alone.

The same restraint applies to structure. An anecdote may create stakes; a detour may reveal judgment. Keep either when it helps. Reorganize only when the current order hides priority, repeats a claim, or forces the reader to reconstruct the logic.

2.4 Each sentence needs a reason

Each sentence should give the reader a claim, condition, example, transition, or moment of voice. This standard is more precise than saying that a sentence must "earn its place." It identifies the sentence's purpose for the reader. Throat-clearing, empty qualifiers, and generic encouragement fail this test because they delay the next useful piece of information.

3 Replace performance with information

Agent-like prose often performs importance before it supplies evidence. Ordinary events become "pivotal moments." A town "boasts" features. A design "underscores" a commitment. These phrases ask the reader to accept an evaluation that the sentence has not supported.

Agent-like draft. The new archive marks a pivotal moment in the institution's evolving landscape, underscoring its enduring commitment to accessible scholarship.

Reader-centered revision. The institution opened the archive to the public and placed the catalog online.

The revision states observable changes. If the date, catalog size, or access policy matters, add it from a source. Do not replace one vague compliment with another.

3.1 Name the relation

Abstract wrappers postpone the claim. Phrases such as "the mechanism behind X," "the role of Y in Z," and "the interplay between A and B" name a relationship without specifying it. Replace the wrapper with the relevant action or cause when the evidence supports one.

Simple verbs often produce the clearest sentence. Prefer "is" and "has" to "serves as," "stands as," or "boasts" when no distinction is lost. Replace "has the ability to process" with "can process." Use a stronger verb only when it names a stronger action.

3.2 Name the source

“Experts argue” and “studies show” borrow authority from unnamed sources. Name the expert, study, dataset, or observation. If no source supports the claim, remove it or state the narrower fact that is available.

Lists of media outlets create a similar problem when they stand in for substance. A citation matters because of the claim it supports, not because several publication names look impressive. Give the relevant context for one source rather than a catalog with no explanation.

3.3 Analysis needs a mechanism

Trailing participles such as “highlighting,” “showcasing,” and “underscoring” often append interpretation without evidence. A real explanation names the mechanism, comparison, or consequence.

Agent-like draft. The interface uses blue and green, reflecting the team’s commitment to trust and growth while fostering a more engaging experience.

Reader-centered revision. The interface uses blue for selected controls and green for successful actions.

The second sentence explains the colors’ function. If the draft does not contain that information, retain only the supported description.

3.4 End on a concrete point

Generic conclusions promise a bright future, celebrate a journey, or announce that a change is a step in the right direction. Such endings add mood without adding information. End with the final supported result, limitation, decision, or next action. A conclusion is useful when it changes the reader’s understanding, not when it merely signals that the document is over.

4 The sentence needs a global focus

A sequence can sound natural word by word while the full sentence points at the wrong object. Editing therefore requires more than replacing suspicious phrases. Follow the subject, verb, and head noun across the entire clause.

4.1 Direct answers avoid question wrappers

Agents often introduce a direct claim with a question word even when no question remains open. Replace the wrapper with a noun or declarative clause when the meaning allows.

Agent-like draft. Each contribution should state what the study learned. The next section explains how the controller filters sensor noise.

Reader-centered revision. Each contribution should state the outcome of the study. The controller filters sensor noise by averaging five readings.

Keep a question-word clause when it represents genuine uncertainty. “We do not know why the controller fails” should remain unresolved until the evidence identifies a cause.

4.2 The head noun stays in focus

Relative clauses can shift attention from the object being described to a new subject. Consider the following edit:

Agent-like draft. Include only setup that a result needs.

Reader-centered revision. Include only the setup that is required by the result.

The bigram “result needs” is locally fluent, but it turns “result” into the subject of the relative clause and makes “setup” its object. The revision keeps “setup” as the grammatical focus. This passive construction is useful because the head noun matters more than the actor. Retain active voice when the new subject identifies responsibility or performs an important action.

4.3 Agency should be visible

Active voice is usually clearer when a person or system acts. Replace “The results are preserved automatically” with “The system preserves the results automatically.” Replace the fragment “No configuration file needed” with “You do not need a configuration file.”

Passive voice remains appropriate when the object or procedure deserves the focus, when the actor is unknown, or when naming the actor would distract from the claim. Choose the construction that puts the needed object in focus and identifies a responsible actor when responsibility matters.

4.4 Packed lists hide relationships

A long comma series can hide the relation among ideas. Keep a complete taxonomy, procedure, checklist, or set of options when the reader needs every member. Otherwise group the items, identify the decisive one, or split the sentence.

Agent-like draft. The report addresses hiring, onboarding, retention, scheduling, purchasing, and reporting.

Reader-centered revision. The report covers staffing and daily operations. Staffing includes hiring, onboarding, and retention; daily operations include scheduling, purchasing, and reporting.

The revision gives the reader two categories rather than six equally weighted terms. It also makes later reference easier.

4.5 One uncertainty marker is enough

Phrases such as “could potentially possibly be argued” stack several weak signals without clarifying the author’s confidence. Choose the one modal or qualification that matches the evidence: “may,” “likely,” “under these conditions,” or “the available data suggest.” More words do not create more precision.

5 Structure should show priority

A document can look organized without helping anyone find the main point. Agent-like answers often add headings and bullets before mapping the dependencies among ideas. The result gives every item the same visual weight.

5.1 Lists need a real set

A list helps when readers need to scan a set or follow a sequence. Procedures and option sets often qualify. Connected claims usually belong in prose because their relation matters more than their count.

A useful editing heuristic is to avoid introducing lists longer than five items unless the complete set matters. This is not a mathematical limit. It forces the writer to distinguish one coherent set from several unnamed groups.

Do not force material into threes. Three claims, adjectives, or examples should appear together because the subject has three relevant parts, not because a triad sounds finished. Avoid false ranges for the same reason. “From hiring to reporting” implies a scale that does not exist; “staffing and reporting” simply names the topics.

5.2 Headings should help navigation

A heading should mark a real change in question or task. Remove headings over one-line sections and sentences that merely repeat the heading. Bold inline labels often make ordinary prose look like a dashboard.

Announcements such as “Let’s dive in,” “Here’s what you need to know,” and “Now we will explore” tell the reader that content is coming. Start with the content. Similarly, documentation should describe the current system rather than narrate the latest diff unless the document is a changelog, release note, or migration guide.

5.3 Rhythm without manufactured drama

Natural prose changes pace. It may use a short sentence for emphasis, but it does not make every line a punchline. Combine stacked fragments when they manufacture urgency:

Agent-like draft. Then the new model arrived. No warning. No compromise. The old rules were gone.

Reader-centered revision. The new model changed the comparison because it did not share the assumptions of the earlier systems.

The revision gives the reader a reason for the change. It does not rely on cinematic timing to make the change feel important.

Avoid slogan formulas for the same reason. “Symmetry is the language of trust” sounds memorable but leaves the relation unspecified. If symmetric layouts make controls easier to predict, state that claim and its conditions.

5.4 Formatting is weak evidence

Mechanical boldface, title case in every heading, decorative emoji, and inline-header lists often accompany agent-like prose. Remove them when they obscure hierarchy or clash with the publication. Do not treat them as proof of authorship.

Punctuation also follows the venue and the writer. The Humanizer pass removes dashes because agents often use them as a default rhythm device. The broader general-writing standard allows an occasional dash in a long piece when it fits the writer’s sample better than a comma, colon, or parenthesis. Curly quotation marks are likewise a typographic convention, not a reliable signal. Match the document’s style rather than changing punctuation for forensic reasons.

5.5 Repeat terms when consistency helps

Synonym cycling makes a stable object appear to change. If a paragraph begins with “the model,” do not rotate through “the system,” “the framework,” and “the approach” merely to avoid repetition. Repeat the clear term until a real distinction requires another one. Controlled repetition reduces the reader’s memory burden.

6 Keep the writer in the prose

Situated attention is a stronger human signal than a typo. The writer notices the detail, tension, or exception that matters to this subject and reader.

6.1 Specific details resist substitution

Names, measurements, dates, odd quotations, and concrete observations resist generic paraphrase. Protect them when the source supports them. Agents often round a specific fact into broad significance; an overzealous editor can make the same mistake while trying to improve flow.

Specificity does not license invention. Do not add a vivid anecdote, opinion, or sensory detail merely to make a passage sound human. The useful detail must belong to the writer or the evidence.

6.2 Tension can remain unresolved

People often hold mixed views. A writer may think a policy is mostly useful and still distrust one consequence. A study may produce a strong result under one condition and an unexplained failure under another. Preserve that tension when the evidence does not settle it.

Agent-like prose tends to close the loop. It converts uncertainty into a neat lesson, adds a positive final paragraph, or balances every criticism with an optimistic future. A human revision can stop at the unresolved fact.

6.3 Defensible choices carry voice

A distinctive word, abrupt transition, or parenthetical aside may look inefficient in isolation. Keep it when the writer can explain its effect on the reader. Deliberate choices create voice. Random irregularity does not.

Era-bound references, local language, and field-specific phrasing may also be part of the voice. Preserve them when the intended audience understands them. Explain them when the audience does not. Do not replace them with generic internet language.

6.4 Polish does not imply automation

Perfect grammar, formal vocabulary, a dry register, or a consistent style can all come from a human author or editor. So can a letter-style greeting, a common transition, a dash, and an isolated short sentence. Unsourced prose is a research problem, not an authorship test. Correct formatting may come from a template.

Secondhand language needs special care. A suspicious phrase inside a quotation, title, proper name, or example belongs to its source. Do not “humanize” quoted material. Explain or cite it as needed while preserving the words.

Version history is stronger evidence than stylistic guesswork. Text committed before the public release of ChatGPT on November 30, 2022, cannot have been written with ChatGPT at that time. Provenance, drafts, and revision records should outweigh a detector’s impression of the final prose.

6.5 Personality depends on the genre

Essays and personal writing can support humor, first person, and open uncertainty. Technical or legal prose may need a neutral voice. A humanizing pass should match the genre instead of adding personality everywhere.

7 Possible causes of agent-like prose

Patterns describe an output; they do not reveal its cause. A deployed language agent may combine pretraining, post-training, system instructions, decoding, retrieval, and tools. Any of these may affect the prose. The ideas below are starting points for experiments, and several may be true at once.

7.1 Local prediction can outrun global structure

During pretraining, a language model learns next-token prediction (NTP). This rewards a continuation that fits the immediate context. Later training changes the model’s behavior, but familiar phrases and smooth word pairs may still win even when the paragraph needs a different construction.

The phrase “result needs” illustrates the possibility. It is locally fluent, yet the full relative clause should keep “setup” in focus. The same pressure might contribute to stock transitions and repeated rhetorical frames. Penalties against repetition could create the opposite surface effect: synonym cycling.

If local prediction is part of the cause, agent prose should make more long-range focus errors than local grammar errors. One test would compare the same model on sentence and paragraph edits, then repeat the paragraph task with a required outline. The comparison should track factual preservation, head-noun focus, and the relation between paragraph openings and conclusions.

7.2 A generic reader produces generic prose

A prompt names a task, but it may not give the model a stable picture of the person using the answer. Generic helpfulness can fill the gap. The prose explains broadly, announces transitions, and calls something “important” without selecting the consequence that matters.

Readerless corporate shorthand fits this account. “Give each section one job” travels easily from one document to another. “Ensure each section helps a reviewer answer one question” depends on a particular reader and task.

This hypothesis predicts that a concrete audience will change the choice of detail. Compare a generic request with one that fixes the reader’s prior knowledge, decision, and reading conditions. Keep the subject and length constant, then ask members of that audience to rate each version’s support for the stated use.

7.3 Mixed training data pulls toward a middle register

Training data spans many genres, institutions, eras, and levels of quality. Without a strong voice cue, a model may settle on language that is acceptable across much of that mixture: common evaluative words, regular paragraphs, and few risky specifics.

Code, API documentation, issue trackers, and technical manuals may add another influence. Those sources favor modular headings, explicit labels, edge cases, and procedural verbs. Carried into an essay, these habits can produce tiny sections and prose that resembles an interface. This is a proposal about transfer from coding data, not a claim that technical documentation makes writing worse.

Style examples offer a simple test. Compare output with no sample, one human sample, and several samples from the same genre. A separate comparison across models with different documented training emphases could test the coding-transfer account.

7.4 Post-training rewards visible helpfulness

Instruction tuning and preference optimization may reward answers that look complete and cooperative. Headings, long lists, recaps, balanced caveats, and offers to continue make effort easy to see. They can be useful in a chat and awkward in a finished document.

These incentives may also reward automatic agreement. “Great question” and “You’re absolutely right” preserve rapport without testing the claim. Stacked hedges can sound cautious while hiding the uncertainty that matters.

Compare a base model and an instruction-tuned model from the same family on the same continuation task. Then ask for a finished document rather than a reply to the user. A drop in chat residue or decorative organization would support this account.

7.5 Pressure to answer invites guessing

Sometimes the prompt and sources do not contain the needed fact, but the agent is still expected to answer. Vague attribution and plausible gap-filling let the prose continue. The model may even turn the absence of evidence into a familiar story.

Retrieval and permission to answer “unknown” should reduce this behavior if completion pressure is responsible. A useful test would ask matched questions with and without retrieval, allow abstention in one condition, and audit every factual claim against the supplied evidence. Fluency should receive no credit in that audit.

These explanations may overlap. Local prediction might supply a familiar phrase, mixed training data might make it generic, and post-training might put it in a helpful-looking list. A weak reader model could then leave irrelevant items in place. An experiment should change one factor at a time and measure a specific language outcome.

8 A five-pass revision

Edit the document in passes so that a style change does not quietly become a factual change.

1. Read the whole draft. Note its central point, intended reader, factual boundary, and a few signs of the writer’s voice.
2. Protect the source material. Mark names, numbers, quotations, citations, technical distinctions, and uncertainty before rewriting.
3. Work at paragraph scale. Find clusters of formulaic language, put the decisive claim at the point needed by the reader, and group details by their relation to that claim.
4. Work at sentence scale. Locate the subject and verb, state direct relations, keep the head noun in focus, and name an actor when agency matters.
5. Compare the result with the source, restore any detail or cadence lost in cleanup, and read the document aloud. Remove chat residue and decorative endings only after the meaning is secure.

8.1 A short worked example

Agent-like draft. Effective writing is not just about clarity; it is about fostering connection, enhancing engagement, and ensuring that ideas resonate. By leveraging a reader-centered approach, writers can create robust communication that truly makes an impact.

Reader-centered revision. Effective writing gives a particular reader the information needed to understand or use an idea.

The first draft piles a canned contrast, a forced group of three, promotional language, and a corporate verb onto a simple idea. It invokes the reader without naming a need. The revision names that need and stops.

This example does not imply that every paragraph should collapse into one sentence. Detail should remain when it supplies evidence, scope, mechanism, or voice. The point is to remove language that performs significance without carrying information.

An editor can remove so much irregularity that the writer disappears. Stop before the document becomes yours. The finished version should still contain the writer's facts, judgment, and cadence. Leave strong sentences alone, even when you would have written them differently.

A Pattern map

This appendix preserves the complete Humanizer catalog in compact form. The numbers match version 2.12.0 of the skill. Each item names a recurring pattern and the corresponding revision. The main chapters explain the shared principles behind the catalog.

A.1 Inflated or unsupported content

1. **Undue emphasis on significance.** Replace claims about pivotal moments, legacies, and broader trends with the event and its supported consequence.
2. **Notability and media lists.** Give the context that makes one citation relevant instead of using publication names as proof of importance.
3. **Superficial *-ing* analysis.** Replace trailing participles such as “highlighting” and “showcasing” with a mechanism, comparison, or result.
4. **Promotional language.** Remove travel-copy adjectives, superlatives, and claims that an object “boasts” a feature. State the observable fact.
5. **Vague attribution.** Name the source behind “experts argue” or “studies show.” Narrow or remove an unsupported claim.
6. **Formulaic challenge and outlook sections.** Replace a generic obstacle-plus-optimism ending with the specific constraint, response, or unresolved problem.

A.2 Lexical and grammatical defaults

7. **Overused agent vocabulary.** Prefer ordinary, precise words to recurrent terms such as “pivotal,” “intricate,” “landscape,” “tapestry,” and “underscore” when those terms add no meaning.
8. **Copula avoidance.** Use “is” and “has” when “serves as,” “stands as,” “features,” or “boasts” only makes the sentence longer.
9. **Negative parallelism and tailing negation.** State the point directly instead of using “not only,” “not just,” or a clipped ending such as “no guessing.”
10. **Rule-of-three overuse.** Keep three-part groups only when the subject has three relevant parts.
11. **Synonym cycling.** Repeat the stable term instead of rotating through near-synonyms for variety.
12. **False ranges.** Use “from X to Y” only when X and Y define a meaningful scale or progression.
13. **Passive voice and subjectless fragments.** Name the actor when agency matters, and turn fragments into sentences when a reader needs the subject.

A.3 Punctuation and visible formatting

14. **Dash-heavy rhythm.** Replace habitual dashes with periods, commas, colons, parentheses, or a new sentence. Match a supplied author sample when dashes are part of that voice.
15. **Mechanical boldface.** Reserve emphasis for a genuine hierarchy or a term that the publication requires.

16. **Inline-header lists.** Turn bold label-and-colon bullets into prose unless readers need to scan independent fields.
17. **Title case everywhere.** Use the heading capitalization required by the venue; sentence case is usually quieter.
18. **Decorative emoji.** Remove icons that decorate headings or bullets without carrying information.
19. **Curly quotation marks as a default tell.** Match the document's typographic convention. Curly marks alone say nothing about authorship.

A.4 Chat residue and missing evidence

20. **Collaborative chat artifacts.** Remove "Of course," "I hope this helps," offers to continue, and other dialogue left inside a finished document.
21. **Knowledge-cutoff disclaimers and gap-filling.** State that a source does not establish a fact, or omit the passage. Do not convert missing evidence into a plausible story.
22. **Sycophantic tone.** Respond to the substance rather than praising the question or agreeing automatically.

A.5 Filler, hedging, and endings

23. **Filler phrases.** Replace "in order to," "due to the fact that," and "has the ability to" with "to," "because," and "can."
24. **Excessive hedging.** Use the one qualification that matches the evidence.
25. **Generic positive conclusions.** End on the last concrete result, limitation, decision, or next action.
26. **Uniform hyphenation.** Hyphenate a compound when the publication and grammatical position require it. Do not copy an attributive hyphen mechanically into predicate position.

A.6 Rhetorical and structural formulas

27. **Persuasive authority tropes.** Remove "the real question," "at its core," and similar ceremony. State the question or claim.
28. **Signposting announcements.** Replace "Let's dive in" and "Now we will explore" with the content.
29. **Fragmented headers.** Remove the one-line warm-up that merely repeats a heading, and merge sections that do not support navigation.
30. **Diff-anchored writing.** Describe the current system unless the document is a changelog, release note, or migration guide.
31. **Manufactured punchlines and staccato drama.** Combine stacked fragments and explain the reason a change matters.
32. **Aphorism formulas.** Replace "X is the language of Y" and similar slogans with the concrete relation.
33. **Conversational rhetorical openers.** Remove fake-candid hooks such as "Honestly?" and give the answer directly.

34. **Question-word clause wrappers.** Replace non-question “what,” “how,” or “why” clauses with a direct noun or statement when the relation is known.
35. **List sprawl.** Group connected ideas in prose. Keep a long list only when it represents a complete set that readers need to scan or count.

A.7 Global focus and reader purpose

36. **Relative-clause subject drift.** Follow the head noun across the full clause. Keep it in focus unless a new subject identifies an important actor or bears responsibility.
37. **Readerless corporate shorthand.** Replace instructions about jobs, ownership, deliverables, or earning a place with the understanding, evidence, decision, or action that the prose supports.

A.8 Additional general-writing checks

The general-writing skill adds several broader editing checks around the Humanizer catalog:

Compare a broad cleanup with the source before accepting it. Restore any lost technical distinction, calibrated uncertainty, specific fact, or recognizable cadence.

Cut canned setups such as binary contrasts, negative lists, throat-clearing, faux-insight claims, colon reveals, and rhetorical questions. Repeated sentence shapes also deserve attention when they make unrelated ideas appear equivalent.

Delete fake-profound metaphors, aphoristic kickers, and summary paragraphs that merely restate the document. Keep an anecdote, detour, fragment, or unusual sequence when it carries context, tension, character, or a needed reading effect.

B Revision checklist

Use this checklist after the substantive edit. A failed item calls for another pass, not a justification for the current sentence.

B.1 Meaning and voice

- ☐ Every factual claim, example, statistic, quotation, and opinion comes from the source or the writer.
- ☐ Calibrated uncertainty and technical distinctions remain intact.
- ☐ Distinctive vocabulary, cadence, humor, tension, and useful asides still sound like the writer.
- ☐ Strong sentences remain unchanged when they already serve the reader.
- ☐ The amount of cutting matches the actual problem.

B.2 Reader and clarity

- ☐ The intended reader and the reader’s use for the document are clear.
- ☐ Every paragraph gives the reader a needed claim, condition, example, transition, or moment of voice.
- ☐ Concrete facts survive, and unsupported abstractions have not gained invented detail.
- ☐ Main subjects and verbs are easy to locate.

- ☐ Head nouns retain the right focus across relative clauses, and responsible actors are named when agency matters.

B.3 Language and structure

- ☐ Importance claims, promotional adjectives, vague attribution, and decorative analysis have been replaced with supported information.
- ☐ Filler, stacked hedges, fake-strong verbs, and abstract wrappers have been removed when a direct phrase preserves the meaning.
- ☐ Question-word wrappers have become direct nouns or statements unless the question is genuinely unresolved.
- ☐ Lists represent real sets, and paragraph structure shows priority rather than surface completeness.
- ☐ Headings support navigation; bold labels, decorative emoji, and chat residue do not substitute for hierarchy.
- ☐ Sentence lengths and paragraph openings vary without stacked punchlines or mechanical symmetry.
- ☐ Stable objects keep stable names instead of cycling through synonyms.
- ☐ The conclusion ends on a supported result, limit, decision, or next action.

B.4 Final read

- ☐ The revision sounds natural when read to a sharp colleague.
- ☐ Quoted language, titles, proper names, and technical terms remain unchanged.
- ☐ Punctuation and capitalization match the writer, genre, and venue.
- ☐ The document contains no claim about agent authorship based only on style.
- ☐ The writer would recognize the facts, judgment, and voice as their own.

C Sources and acknowledgments

This guide distills two agent-facing resources in the [writing-skills](#) repository: Humanizer version 2.12.0 and the general-writing skill.

Humanizer builds on Siqi Chen’s [humanizer](#), which draws its pattern catalog from Wikipedia’s [Signs of AI writing](#). The general-writing framework builds on Peter Yang’s [no-ai-slop](#). Both upstream projects are distributed under the MIT License, and their license notices remain with the corresponding skills.

The causal hypotheses in [Section 7](#) are conjectures presented for discussion and testing. They should not be attributed to the upstream projects.