

Human vs. Agent Writing

Max Simchowitz

An editing guide for writers and editors who want to remove agent-like habits without erasing meaning or voice. Use the patterns to diagnose and revise prose, not to identify its author.

Date: July 25, 2026

Abstract

This guide is for writers and editors reviewing prose that feels polished but generic or detached from a particular reader. It treats recurring agent-like habits as editing problems, not evidence about authorship. Use the patterns to locate places where a draft inflates a claim or organizes material mechanically. The same review can expose vague assumptions about the audience. Revisions should preserve supported facts and uncertainty while leaving the writer's deliberate voice intact. The examples and checklist work as a line-editing aid or a final review before publication.

Contents

1	Use patterns to revise prose	3
2	Preserve meaning and voice	3
3	Replace vague emphasis with facts	4
4	Read the whole sentence	5
5	Use structure to show priority	7
6	Preserve the writer's voice	8
7	Possible causes of recurring agent-like patterns	9
8	Revise in five passes	10
A	Pattern map	12
B	Revision checklist	14
C	Sources and acknowledgments	15

1 Use patterns to revise prose

People write into situations. They anticipate objections and select facts according to what this audience already knows. They also keep uncertainty that cannot honestly be removed. An agent can imitate these choices. Without a clear audience, it often produces a generic answer.

That is the distinction this guide cares about. Grammar is weak evidence about the source of a passage; formal vocabulary and punctuation do not add much. Human writers use formulas, while agents can write restrained, specific prose. A clean sentence proves nothing about authorship.

1.1 Identify the intended reader

Start an edit by naming the reader and the consequence of the document. A reviewer checking evidence needs different details from a student learning a mechanism. A collaborator choosing between two methods needs the comparison early. Once that purpose is clear, the editor can judge each paragraph by the new understanding or decision it gives that reader.

This test catches vague management language. “Give each section one job” sounds firm but explains little. “Ensure each section helps the reviewer answer one question” gives the rule a reason.

1.2 Make claims about the passage

Describe the passage precisely. Call it *formulaic* or *poorly sourced*. If it overstates the evidence, say so. Each of those claims can be checked. A claim about the person or system that wrote the passage needs independent evidence.

Style-based detection produces familiar false positives. Professional editing can explain consistent grammar. Word processors explain curly quotation marks, while a field’s shared vocabulary may explain repeated terms. A journalist may use many dashes. Look for a cluster of choices that causes a reading problem rather than a lone tell.

2 Preserve meaning and voice

A successful edit changes the prose without changing the writer’s claim. Style is subordinate to factual fidelity. Names and numbers in the revision must come from the source or the writer. So must dates and quotations. Preserve citations and technical distinctions under the same rule. When the draft lacks a needed fact, either ask for it or write the honest, less specific sentence.

Fiction is the exception because invention belongs to the task. In factual prose, a vivid unsupported detail remains a fabrication. The same constraint applies to technical, legal, and reference writing.

2.1 Preserve uncertainty

Preserve the difference between “demonstrates,” “suggests,” “is consistent with,” and “does not rule out.” A shorter sentence is worse when it turns a conditional result into a universal one. The same principle applies to opinions. Preserve the writer’s confidence. Hesitation and mixed responses should survive the edit too.

Missing evidence calls for a boundary, not a plausible substitute. If the available sources do not establish a founding date, say so when the absence matters or omit the sentence. Do not infer a private life from silence. Do not invent a motivation or causal history either.

2.2 Preserve deliberate choices

Voice appears in selection as much as style. A writer may keep a technical qualification that interrupts the rhythm because it is true. Another may use a dated joke or pause to correct an earlier thought. Do not regularize those choices merely to make every paragraph sound alike.

When the writer supplies an earlier sample, use it as the strongest evidence for vocabulary and punctuation. Match its rhythm too. A generic style rule should not override a consistent choice in that sample unless the choice harms meaning or violates the publication's requirements.

Voice is often uneven. A short sentence may follow a long one. An aside may reveal the writer's stake, and a passage may remain ambivalent. Keep that unevenness when it belongs to the writer.

2.3 Make the lightest effective edit

Read the full draft before touching a sentence. Identify its central point and intended reader. Then notice a few voice signals worth preserving. Fix the actual problems and leave strong sentences alone.

The same restraint applies to structure. An anecdote may create stakes; a detour may reveal judgment. Keep either when it helps. Reorganize only when the current order hides priority. Repetition may also justify a move, as may an order that forces the reader to reconstruct the logic.

2.4 Check each sentence's purpose

Each sentence should give the reader a claim, condition, example, transition, or moment of voice. This standard is more precise than saying that a sentence must "earn its place." It identifies the sentence's purpose for the reader. Throat-clearing and empty qualifiers fail this test. Generic encouragement does too because it delays the next useful piece of information.

3 Replace vague emphasis with facts

Agent-like prose often performs importance before it supplies evidence. Ordinary events become "pivotal moments." A town "boasts" features. A design "underscores" a commitment. These phrases ask the reader to accept an evaluation that the sentence has not supported.

Agent-like draft

The new archive marks a pivotal moment in the institution's evolving landscape, underscoring its enduring commitment to accessible scholarship.

Reader-centered revision

The institution opened the archive to the public and placed the catalog online.

The revision states observable changes. If the date, catalog size, or access policy matters, add it from a source. Do not replace one vague compliment with another.

3.1 Explain the relation

Abstract wrappers postpone the claim. Phrases such as "the mechanism behind X," "the role of Y in Z," and "the interplay between A and B" name a relationship without specifying it. Replace the wrapper with the relevant action or cause when the evidence supports one.

Simple verbs often produce the clearest sentence. Prefer "is" and "has" to "serves as," "stands as," or "boasts" when no distinction is lost. Replace "has the ability to process" with "can process." Use a stronger verb only when it names a stronger action.

3.2 Name the source

“Experts argue” and “studies show” borrow authority from unnamed sources. Name the expert or study. If the claim comes from a dataset or observation, identify that source instead. If no source supports the claim, remove it or state the narrower fact that is available.

Lists of media outlets create a similar problem when they stand in for substance. A citation matters because of the claim it supports, not because several publication names look impressive. Give the relevant context for one source rather than a catalog with no explanation.

3.3 Explain the mechanism

Trailing participles such as “highlighting,” “showcasing,” and “underscoring” often append interpretation without evidence. A real explanation names the mechanism. It may instead establish a comparison or consequence.

Agent-like draft

The interface uses blue and green, reflecting the team’s commitment to trust and growth while fostering a more engaging experience.

Reader-centered revision

The interface uses blue for selected controls and green for successful actions.

The second sentence explains the colors’ function. If the draft does not contain that information, retain only the supported description.

3.4 End on a concrete point

Generic conclusions may promise a bright future or celebrate a journey. Others announce that a change is a step in the right direction. Such endings add mood without adding information. End with the final supported result. If a limitation controls its interpretation, name that limit. Documents that ask for action should close on the decision or next step. A conclusion is useful when it changes the reader’s understanding, not when it merely signals that the document is over.

4 Read the whole sentence

A local phrase can sound natural while the sentence as a whole feels stiff or describes the wrong relation. Editing requires more than swapping suspicious phrases. Read the subject, verb, and modifiers together, then recast the sentence if the repair sounds mechanical.

4.1 Remove unnecessary question wrappers

Agents often introduce a direct claim with a question word even when no question remains open. Replace the wrapper with a noun or declarative clause when the meaning allows.

Agent-like draft

Each contribution should state what the study learned. The next section explains how the controller filters sensor noise.

Reader-centered revision

Each contribution should state the outcome of the study. The controller filters sensor noise by averaging five readings.

Keep a question-word clause when it represents genuine uncertainty. “We do not know why the controller fails” should remain unresolved until the evidence identifies a cause.

4.2 Rewrite mechanical relative clauses

Relative clauses can be grammatical and still sound stiff. Consider the following edit:

Agent-like draft

Include only the setup that is required by the result.

Reader-centered revision

Include only the setup needed to interpret the result.

The first sentence forces a passive relative clause and treats the result as if it requires setup. The revision states why the setup is present. “Give readers only the setup they need to interpret the result” also works when the reader should be explicit. The right revision depends on the sentence’s purpose, not on keeping one noun in a preferred grammatical position.

4.3 Name the actor when it matters

Active voice is usually clearer when a person or system acts. Replace “The results are preserved automatically” with “The system preserves the results automatically.” Replace the fragment “No configuration file needed” with “You do not need a configuration file.”

Passive voice remains appropriate when the object or procedure deserves the focus. It also works when the actor is unknown. If naming the actor would distract from the claim, leave the actor out. Choose the construction that puts the needed object in focus and identifies a responsible actor when responsibility matters.

4.4 Packed lists hide relationships

A comma series can hide the relation among ideas. Length is not the only problem: even three objects can feel arbitrary when they serve different roles. Keep a fixed taxonomy or checklist when the reader needs every member. Procedures and option sets also belong in full.

In ordinary prose, identify the governing relation before retaining a series. If a paragraph contains several inventory sentences, let one state the main set and rewrite the others around categories or actions. A paragraph break alone does not vary the syntax.

Agent-like draft

The report addresses hiring, onboarding, retention, scheduling, purchasing, and reporting.

Reader-centered revision

The report separates staffing from daily operations. Hiring and onboarding fall under staffing; retention does too. Scheduling and purchasing are daily operations, as is reporting.

The revision gives the reader two categories rather than six equally weighted terms. It also varies the sentence structure and makes later reference easier.

4.5 One uncertainty marker is enough

Phrases such as “could potentially possibly be argued” stack several weak signals without clarifying the author’s confidence. Choose the one modal or qualification that matches the evidence: “may,” “likely,” “under these conditions,” or “the available data suggest.” More words do not create more precision.

5 Use structure to show priority

A document can look organized without helping anyone find the main point. Agent-like answers often add headings and bullets before mapping the dependencies among ideas. The result gives every item the same visual weight.

5.1 Use lists for coherent sets

A list helps when readers need to scan a set or follow a sequence. Procedures often qualify; fixed option sets do too. Connected claims usually belong in prose because their relation matters more than their count.

A useful editing heuristic is to avoid introducing lists longer than five items unless the complete set matters. This is not a mathematical limit. It forces the writer to distinguish one coherent set from several unnamed groups.

Do not force material into threes. Three claims should appear together because the subject has three relevant parts, not because a triad sounds finished. The same rule applies to adjectives and examples. Avoid false ranges for the same reason. “From hiring to reporting” implies a scale that does not exist; “staffing and reporting” simply names the topics.

5.2 Use headings for navigation

A heading should mark a real change in question or task. Remove headings over one-line sections and sentences that merely repeat the heading. Bold inline labels often make ordinary prose look like a dashboard.

Announcements such as “Let’s dive in,” “Here’s what you need to know,” and “Now we will explore” tell the reader that content is coming. Start with the content. Similarly, documentation should describe the current system rather than narrate the latest diff unless the document is a changelog, release note, or migration guide.

5.3 Vary rhythm without forcing drama

Natural prose changes pace. It may use a short sentence for emphasis, but it does not make every line a punchline. Combine stacked fragments when they manufacture urgency:

Agent-like draft

Then the new model arrived. No warning. No compromise. The old rules were gone.

Reader-centered revision

The new model changed the comparison because it did not share the assumptions of the earlier systems.

The revision gives the reader a reason for the change. It does not rely on cinematic timing to make the change feel important.

Avoid slogan formulas for the same reason. “Symmetry is the language of trust” sounds memorable but leaves the relation unspecified. If symmetric layouts make controls easier to predict, state that claim and its conditions.

5.4 Formatting is weak evidence

Mechanical boldface and title case in every heading often accompany agent-like prose. Decorative emoji and inline-header lists do too. Remove these devices when they obscure hierarchy or clash with the publication. Do not treat them as proof of authorship.

Punctuation also follows the venue and the writer. The Humanizer pass removes dashes because agents often use them as a default rhythm device. The broader general-writing standard allows an occasional dash in a long piece when it fits the writer's sample better than the alternatives. A comma may work, or the sentence may need a colon or parenthesis. Curly quotation marks are likewise a typographic convention, not a reliable signal. Match the document's style rather than changing punctuation for forensic reasons.

5.5 Repeat stable terms

Synonym cycling makes a stable object appear to change. If a paragraph begins with "the model," do not rotate through "the system," "the framework," and "the approach" merely to avoid repetition. Repeat the clear term until a real distinction requires another one. Controlled repetition reduces the reader's memory burden.

6 Preserve the writer's voice

Situated attention is a stronger human signal than a typo. The writer notices the detail, tension, or exception that matters to this subject and reader.

6.1 Preserve supported specifics

Names and measurements resist generic paraphrase. Dates resist it too, as do odd quotations and concrete observations. Protect these details when the source supports them. Agents often round a specific fact into broad significance; an overzealous editor can make the same mistake while trying to improve flow.

Specificity does not license invention. Do not add a vivid anecdote or opinion merely to make a passage sound human. A sensory detail needs the same support. The useful detail must belong to the writer or the evidence.

6.2 Leave unresolved tension intact

People often hold mixed views. A writer may think a policy is mostly useful and still distrust one consequence. A study may produce a strong result under one condition and an unexplained failure under another. Preserve that tension when the evidence does not settle it.

Agent-like prose tends to close the loop. It converts uncertainty into a neat lesson, adds a positive final paragraph, or balances every criticism with an optimistic future. A human revision can stop at the unresolved fact.

6.3 Keep deliberate choices

A distinctive word or abrupt transition may look inefficient in isolation. So may a parenthetical aside. Keep it when the writer can explain its effect on the reader. Deliberate choices create voice. Random irregularity does not.

Era-bound references and local language may also be part of the voice. Field-specific phrasing can be as well. Preserve these choices when the intended audience understands them. Explain them when the audience does not. Do not replace them with generic internet language.

6.4 Polish does not imply automation

Perfect grammar can come from a human author or editor. So can formal vocabulary or a dry register. A consistent style proves no more. The same is true of a letter-style greeting or common transition. Neither one dash nor an isolated short sentence establishes automation. Unsourced prose is a research problem, not an authorship test. Correct formatting may come from a template.

Secondhand language needs special care. A suspicious phrase inside a quotation or title belongs to its source. The same protection applies to proper names and examples. Do not “humanize” quoted material. Explain or cite it as needed while preserving the words.

Version history is stronger evidence than stylistic guesswork. Text committed before the public release of ChatGPT on November 30, 2022, cannot have been written with ChatGPT at that time. Provenance should outweigh a detector’s impression of the final prose. Drafts and revision records make that provenance easier to establish.

6.5 Match personality to the genre

Essays and personal writing can support humor and first person. They may also leave uncertainty open. Technical or legal prose may need a neutral voice. A humanizing pass should match the genre instead of adding personality everywhere.

7 Possible causes of recurring agent-like patterns

Patterns describe an output; they do not reveal its cause. A deployed language agent begins with pre-training and may later undergo post-training. System instructions and decoding also shape its response. Retrieval or tool use can change the available evidence. Any of these may affect the prose. The ideas below are starting points for experiments, and several may be true at once.

7.1 Local prediction may weaken larger structures

During pretraining, a language model learns next-token prediction (NTP). This rewards a continuation that fits the immediate context. Later training changes the model’s behavior, but familiar phrases and smooth word pairs may still win even when a sentence or paragraph needs a different construction.

Locally plausible continuations can still produce a sentence whose relation is hard to recover. The same pressure might contribute to stock transitions and repeated rhetorical frames. Penalties against repetition could create the opposite surface effect: synonym cycling.

If local prediction is part of the cause, agent prose should make more sentence-level relation errors and paragraph-level focus errors than local grammar errors. One test would compare the same model on sentence and paragraph edits, then repeat the paragraph task with a required outline. The comparison should track factual preservation and whether each edited sentence states its relation directly. It should also compare paragraph openings with their conclusions.

7.2 An underspecified audience may produce generic prose

A prompt names a task, but it may not give the model a stable picture of the person using the answer. Generic helpfulness can fill the gap. The prose explains broadly, announces transitions, and calls something “important” without selecting the consequence that matters.

Readerless corporate shorthand fits this account. “Give each section one job” travels easily from one document to another. “Ensure each section helps a reviewer answer one question” depends on a particular reader and task.

This hypothesis predicts that a concrete audience will change the choice of detail. Compare a generic request with one that states the reader’s prior knowledge and pending decision. The second prompt should also specify the reading conditions. Keep the subject and length constant, then ask members of that audience to rate each version’s support for the stated use.

7.3 Mixed training data may favor a middle register

Training data spans many genres and institutions. It also crosses eras and levels of quality. Without a strong voice cue, a model may settle on language that is acceptable across much of that mixture. The result may rely on common evaluative words and regular paragraphs while avoiding risky specifics.

Code and API documentation may add another influence. Issue trackers and technical manuals may contribute as well. Those sources favor modular headings and explicit labels. They also isolate edge cases and use procedural verbs. Carried into an essay, these habits can produce tiny sections and prose that resembles an interface. This is a proposal about transfer from coding data, not a claim that technical documentation makes writing worse.

Style examples offer a simple test. Compare output with no sample, one human sample, and several samples from the same genre. A separate comparison across models with different documented training emphases could test the coding-transfer account.

7.4 Post-training may reward visible helpfulness

Instruction tuning and preference optimization may reward answers that look complete and cooperative. Headings and long lists make effort easy to see. Recaps and balanced caveats can do the same, as can an offer to continue. These devices can be useful in a chat and awkward in a finished document.

These incentives may also reward automatic agreement. “Great question” and “You’re absolutely right” preserve rapport without testing the claim. Stacked hedges can sound cautious while hiding the uncertainty that matters.

Compare a base model and an instruction-tuned model from the same family on the same continuation task. Then ask for a finished document rather than a reply to the user. A drop in chat residue or decorative organization would support this account.

7.5 Answer pressure may increase guessing

Sometimes the prompt and sources do not contain the needed fact, but the agent is still expected to answer. Vague attribution and plausible gap-filling let the prose continue. The model may even turn the absence of evidence into a familiar story.

Retrieval and permission to answer “unknown” should reduce this behavior if completion pressure is responsible. A useful test would ask matched questions with and without retrieval. One condition should allow abstention. The final audit should check every factual claim against the supplied evidence, with no credit for fluency.

These explanations may overlap. Local prediction might supply a familiar phrase. Mixed training data could make the phrase generic, while post-training could place it in a helpful-looking list. A weak reader model might then leave irrelevant items in place. An experiment should change one factor at a time and measure a specific language outcome.

8 Revise in five passes

Edit the document in passes so that a style change does not quietly become a factual change.

1. Read the whole draft. Note its central point and intended reader. Record the factual boundary and a few signs of the writer's voice.
2. Protect the source material. Mark names and numbers first. Protect quotations and citations too. Preserve technical distinctions and uncertainty before rewriting.
3. Work at paragraph scale. Find clusters of formulaic language. Put the decisive claim at the point needed by the reader, then group details by their relation to that claim.
4. Work at sentence scale. Locate the subject and verb, then check whether each modifier states its relation directly. Name an actor when agency matters.
5. Compare the result with the source. Restore any detail or cadence lost in cleanup, then read the document aloud. Remove chat residue and decorative endings only after the meaning is secure.

8.1 Worked example

Agent-like draft

Effective writing is not just about clarity; it is about fostering connection, enhancing engagement, and ensuring that ideas resonate. By leveraging a reader-centered approach, writers can create robust communication that truly makes an impact.

Reader-centered revision

Effective writing gives a particular reader the information needed to understand or use an idea.

The first draft puts a canned contrast beside a forced group of three. Promotional language and a corporate verb add ceremony to the same simple idea. The draft invokes the reader without naming a need. The revision names that need and stops.

This example does not imply that every paragraph should collapse into one sentence. Detail should remain when it supplies evidence or scope. A mechanism or moment of voice may also be worth keeping. The point is to remove language that performs significance without carrying information.

An editor can remove so much irregularity that the writer disappears. Stop before the document becomes yours. The finished version should still contain the writer's facts and judgment. Its cadence should remain recognizable too. Leave strong sentences alone, even when you would have written them differently.

A Pattern map

This appendix preserves the complete Humanizer catalog in compact form. The numbers match version 2.13.0 of the skill. Each item names a recurring pattern and the corresponding revision. The main chapters explain the shared principles behind the catalog.

A.1 Inflated or unsupported content

1. **Undue emphasis on significance.** Replace claims about pivotal moments, legacies, and broader trends with the event and its supported consequence.
2. **Notability and media lists.** Give the context that makes one citation relevant instead of using publication names as proof of importance.
3. **Superficial *-ing* analysis.** Replace trailing participles such as “highlighting” and “showcasing” with a mechanism, comparison, or result.
4. **Promotional language.** Remove travel-copy adjectives, superlatives, and claims that an object “boasts” a feature. State the observable fact.
5. **Vague attribution.** Name the source behind “experts argue” or “studies show.” Narrow or remove an unsupported claim.
6. **Formulaic challenge and outlook sections.** Replace a generic obstacle-plus-optimism ending with the specific constraint, response, or unresolved problem.

A.2 Lexical and grammatical defaults

7. **Overused agent vocabulary.** Prefer ordinary, precise words to recurrent terms such as “pivotal,” “intricate,” “landscape,” “tapestry,” and “underscore” when those terms add no meaning.
8. **Copula avoidance.** Use “is” and “has” when “serves as,” “stands as,” “features,” or “boasts” only makes the sentence longer.
9. **Negative parallelism and tailing negation.** State the point directly instead of using “not only,” “not just,” or a clipped ending such as “no guessing.”
10. **Rule-of-three overuse.** Keep three-part groups only when the subject has three relevant parts.
11. **Synonym cycling.** Repeat the stable term instead of rotating through near-synonyms for variety.
12. **False ranges.** Use “from X to Y” only when X and Y define a meaningful scale or progression.
13. **Passive voice and subjectless fragments.** Name the actor when agency matters, and turn fragments into sentences when a reader needs the subject.

A.3 Punctuation and visible formatting

14. **Dash-heavy rhythm.** Replace habitual dashes with periods, commas, colons, parentheses, or a new sentence. Match a supplied author sample when dashes are part of that voice.
15. **Mechanical boldface.** Reserve emphasis for a genuine hierarchy or a term that the publication requires.

16. **Inline-header lists.** Turn bold label-and-colon bullets into prose unless readers need to scan independent fields.
17. **Title case everywhere.** Use the heading capitalization required by the venue; sentence case is usually quieter.
18. **Decorative emoji.** Remove icons that decorate headings or bullets without carrying information.
19. **Curly quotation marks as a default tell.** Match the document's typographic convention. Curly marks alone say nothing about authorship.

A.4 Chat residue and missing evidence

20. **Collaborative chat artifacts.** Remove "Of course," "I hope this helps," offers to continue, and other dialogue left inside a finished document.
21. **Knowledge-cutoff disclaimers and gap-filling.** State that a source does not establish a fact, or omit the passage. Do not convert missing evidence into a plausible story.
22. **Sycophantic tone.** Respond to the substance rather than praising the question or agreeing automatically.

A.5 Filler, hedging, and endings

23. **Filler phrases.** Replace "in order to," "due to the fact that," and "has the ability to" with "to," "because," and "can."
24. **Excessive hedging.** Use the one qualification that matches the evidence.
25. **Generic positive conclusions.** End on the last concrete result, limitation, decision, or next action.
26. **Uniform hyphenation.** Hyphenate a compound when the publication and grammatical position require it. Do not copy an attributive hyphen mechanically into predicate position.

A.6 Rhetorical and structural formulas

27. **Persuasive authority tropes.** Remove "the real question," "at its core," and similar ceremony. State the question or claim.
28. **Signposting announcements.** Replace "Let's dive in" and "Now we will explore" with the content.
29. **Generic, fragmented, or overbuilt headings.** Review titles, section and subsection headings, and captions as prose. Replace vague or mechanical display text, remove warm-ups that repeat a heading, and merge sections that do not support navigation.
30. **Diff-anchored writing.** Describe the current system unless the document is a changelog, release note, or migration guide.
31. **Manufactured punchlines and staccato drama.** Combine stacked fragments and explain the reason a change matters.
32. **Aphorism formulas.** Replace "X is the language of Y" and similar slogans with the concrete relation.
33. **Conversational rhetorical openers.** Remove fake-candid hooks such as "Honestly?" and give the answer directly.

34. **Question-word clause wrappers.** Replace non-question “what,” “how,” or “why” clauses with a direct noun or statement when the relation is known.
35. **Loose catalogs and list sprawl.** Require a governing relation for every series. Rewrite adjacent inventory sentences instead of hiding the same rhythm behind paragraph breaks. Keep a long list only when readers need to scan or count a fixed set.

A.7 Sentence focus and reader purpose

36. **Mechanical relative clauses.** Read the full sentence and reduce or recast a stiff relative clause. Do not force active or passive voice merely to keep one noun in a preferred grammatical position.
37. **Readerless corporate shorthand.** Replace instructions about jobs, ownership, deliverables, or earning a place with the understanding, evidence, decision, or action that the prose supports.

A.8 Additional general-writing checks

The general-writing skill adds several broader editing checks around the Humanizer catalog:

Compare a broad cleanup with the source before accepting it. Restore any lost technical distinction or calibrated uncertainty. Keep specific facts and recognizable cadence.

Cut binary contrasts and negative lists when a direct claim says the same thing. Throat-clearing and faux-insight claims also delay the point. Rewrite colon reveals or rhetorical questions that manufacture drama. Repeated sentence shapes deserve attention when they make unrelated ideas appear equivalent.

Delete fake-profound metaphors and aphoristic kickers. Remove summary paragraphs that merely restate the document. Keep an anecdote or detour when it carries context or tension. A fragment or unusual sequence can also remain when it adds character or creates a needed reading effect.

B Revision checklist

Use this checklist after the substantive edit. A failed item calls for another pass, not a justification for the current sentence.

B.1 Meaning and voice

- ☐ Every factual claim and example comes from the source or the writer. Statistics, quotations, and opinions do too.
- ☐ Calibrated uncertainty and technical distinctions remain intact.
- ☐ Distinctive vocabulary and cadence still sound like the writer. Preserve the writer’s humor, tension, and useful asides as well.
- ☐ Strong sentences remain unchanged when they already serve the reader.
- ☐ The amount of cutting matches the actual problem.

B.2 Reader and clarity

- ☐ The intended reader and the reader’s use for the document are clear.
- ☐ Every paragraph gives the reader a needed claim, condition, example, transition, or moment of voice.

- ☐ Concrete facts survive, and unsupported abstractions have not gained invented detail.
- ☐ Main subjects and verbs are easy to locate.
- ☐ Relative clauses and passive constructions sound natural in the full sentence. Responsible actors are named when agency matters.

B.3 Language and structure

- ☐ Unsupported importance claims and promotional adjectives have become sourced facts. Treat vague attribution and decorative analysis the same way.
- ☐ Filler and stacked hedges have been removed when a direct phrase preserves the meaning. Fake-strong verbs and abstract wrappers have too.
- ☐ Question-word wrappers have become direct nouns or statements unless the question is genuinely unresolved.
- ☐ Every series has a governing relation. Ordinary paragraphs do not stack list-shaped sentences or hide repetition behind paragraph breaks.
- ☐ The title and every section or subsection heading accurately frame the material that follows.
- ☐ Captions and example labels state the evidence or action they frame. Bold labels do not substitute for hierarchy.
- ☐ Sentence lengths and paragraph openings vary without stacked punchlines or mechanical symmetry.
- ☐ Stable objects keep stable names instead of cycling through synonyms.
- ☐ The conclusion ends on a supported result or limit. When the document calls for action, it ends on the decision or next step.

B.4 Final read

- ☐ The revision sounds natural when read to a sharp colleague.
- ☐ Quoted language and titles remain unchanged. Proper names and technical terms do too.
- ☐ Punctuation and capitalization match the writer, genre, and venue.
- ☐ The document contains no claim about agent authorship based only on style.
- ☐ The writer would recognize the facts and judgment as their own. The voice remains recognizable too.

C Sources and acknowledgments

This guide distills two agent-facing resources in the [writing-skills](#) repository: Humanizer version 2.13.0 and the general-writing skill.

Humanizer builds on Siqi Chen’s [humanizer](#), which draws its pattern catalog from Wikipedia’s [Signs of AI writing](#). The general-writing framework builds on Peter Yang’s [no-ai-slop](#). Both upstream projects are distributed under the MIT License, and their license notices remain with the corresponding skills.

The causal hypotheses in [Section 7](#) are conjectures presented for discussion and testing. They should not be attributed to the upstream projects.