

Human and Agent Writing at the Passage Level

Max Simchowitz

Why fluent prose can still feel generic. Surface fluency can conceal failures of emphasis, agency, and scope. A passage-level analysis can identify these failures without turning style into a verdict about authorship.

Date: July 26, 2026

Abstract

Fluent prose can feel generic even when every sentence is grammatical. When no reader is specified, agent-generated text can give routine background and decisive evidence equal weight, then recast unresolved uncertainty in familiar formulas. Human writers also use these formulas, and language agents can produce precise, situated prose, so surface patterns cannot identify authorship. These patterns are better understood as properties of passages and as clues that motivate hypotheses about the pressures shaping model output. Revision therefore begins by identifying the intended reader and recovering the relations among a claim, its evidence, and the conditions under which it holds. Explanations based on audience specification, training, or local prediction remain hypotheses. A successful revision makes those relations easier to understand without erasing the writer's judgment or changing the strength of the claim.

Contents

1	The passage is the unit of comparison	3
2	Fidelity sets the limits of revision	3
3	Generic emphasis competes with evidence	4
4	Local fluency can conceal relational errors	6
5	Visible organization can flatten priority	7
6	Voice emerges from situated attention	8
7	Possible causes of recurring agent-like patterns	10
8	Revision as constrained inference	11
A	Pattern taxonomy	13
B	Dimensions of comparative analysis	15
C	Sources and acknowledgments	17

1 The passage is the unit of comparison

I use *situated* to describe prose shaped for a particular reader and purpose. Its author anticipates likely objections, selects facts according to what that reader already knows, and preserves uncertainty that the evidence cannot resolve. Language agents can imitate every one of these choices. When a prompt leaves the situation underspecified, however, their answers can converge on a portable middle register broad enough for many readers without being fully suited to any of them.

The pattern catalog is descriptive: it records recurrent editorial problems but does not estimate their frequency. Comparisons based on the catalog are therefore probabilistic and properly made at the passage level. Grammar, formal vocabulary, and conventional punctuation provide weak evidence about provenance because each has ordinary explanations. A clean sentence establishes little about its author.

1.1 Audience specification changes the choice of detail

The intended reader determines which background matters and when a comparison becomes useful. A reviewer checking the evidence may look first for experimental controls; a student learning the mechanism may need conceptual setup before those controls have meaning. A collaborator choosing between methods faces a different question and needs the governing tradeoff early. The technical subject can remain the same even though the reader's purpose changes the explanation.

Generic management language obscures this dependence. "Give each section one job" works across many contexts because it does not identify a reading situation. "Each section should help a reviewer answer one question" is narrower: it connects the section's organization to a particular act of evaluation. The review task determines the structure; merely adding the word *reader* would not make the sentence specific.

1.2 Pattern claims remain properties of passages

Stylistic labels are reliable only when they describe properties that can be checked against the passage. A statement about the person or system that produced it requires independent evidence. This distinction matters because style-based detection has familiar false positives: professional editing explains consistent grammar, word processors explain curly quotation marks, and a field's shared vocabulary explains many repeated terms.

Clusters still matter, but for diagnosis rather than attribution. Inflated emphasis becomes more damaging when it coincides with generic attribution and a paragraph whose sentences all share one shape. Together, those choices can hide the claim and its support, yet their evidential value remains at the passage level.

2 Fidelity sets the limits of revision

A successful revision changes the reading experience without changing the writer's claim. Style remains subordinate to factual fidelity: names, numbers, dates, and quotations remain grounded in the source or the writer. The same constraint protects citations and technical distinctions. When the draft lacks a needed fact, vivid detail cannot fill the gap.

Fiction differs because invention belongs to the form. In factual prose, an unsupported sensory detail remains a fabrication no matter how natural it sounds. Technical, legal, and reference writing impose the same boundary.

2.1 Uncertainty is part of the claim

The verbs “demonstrates,” “suggests,” “is consistent with,” and “does not rule out” encode different relations between evidence and conclusion. Compression becomes distortion when it turns a conditional result into a universal one. Confidence and ambivalence require similar protection: a writer’s hesitation may be the most accurate account of the evidence.

Missing evidence also constrains what the passage can say. If the available sources do not establish a founding date, the passage can acknowledge that absence when it matters or leave the date unmentioned. A plausible biography, motivation, or causal history would fill the gap with material that the sources do not support.

2.2 The choice of detail carries voice

Voice appears in the choice of detail as much as in punctuation. One writer may retain a technical qualification that interrupts the rhythm because the qualification prevents an overclaim. Another may pause for a dated joke or correct an earlier thought in a parenthesis. Regularizing either choice can make the prose smoother and less faithful at the same time.

An earlier sample from the same writer provides stronger evidence about diction and cadence than a generic style rule. Such samples also reveal that voice is often uneven: a long explanation may give way to a short conclusion, while an aside interrupts an otherwise formal register. That unevenness can be part of the voice, especially when the writer leaves a genuine tension unresolved.

2.3 The smallest sufficient intervention

The scale of revision follows the scale of the defect. A confused modifier may need a local repair, whereas a paragraph that hides its central comparison may need a new order. Anecdotes and detours survive when they create stakes or reveal judgment. Reorganization becomes warranted when the existing sequence obscures priority, repeats the same point, or forces the reader to reconstruct a dependency.

Limiting the edit in this way leaves strong sentences intact and prevents the editor’s preferred cadence from replacing the writer’s merely because both are defensible.

2.4 Sentence function depends on the reader

A sentence can contribute a claim, condition, example, transition, or moment of voice. Its value lies in the understanding or inference it makes possible, not in an abstract requirement to “earn its place.” Throat-clearing and empty qualifiers fail because they postpone that contribution. Generic encouragement has the same effect when it replaces information the reader needs.

3 Generic emphasis competes with evidence

Agent-like prose can declare a subject important before giving the reader evidence for that judgment. Ordinary events become “pivotal moments,” a town “boasts” features, and a design “underscores” a commitment. Such phrases ask the reader to accept an evaluation that the sentence has not yet supported.

Agent-like draft

The new archive marks a pivotal moment in the institution’s evolving landscape, underscoring its enduring commitment to accessible scholarship.

Reader-centered revision

The institution opened the archive to the public and placed the catalog online.

The revision replaces an evaluation with two observable changes. A date, catalog size, or access policy could make the account more useful if the source supplied it; absent that evidence, another compliment would only restate the original problem.

3.1 Direct verbs expose the relation

Abstract wrappers postpone the action. Phrases such as “the mechanism behind X” or “the role of Y in Z” identify a relationship without specifying it. A direct verb reveals whether X causes, constrains, predicts, or merely correlates with Y, provided the evidence supports that relation.

Verb specificity also depends on the subject. A claim can hold under stated conditions and imply consequences; motion toward a stopping point describes neither relation. Conditions do not simply “limit the claim,” either, because that wording leaves the affected dimension unnamed. A validity statement names the conditions under which a claim holds, whereas an implication statement specifies how far its consequences extend. Sometimes the sentence needs an intermediate noun such as *extent*, *applicability*, or *consequence* before the verb can express the relation accurately.

Simple copulas are often sufficient. “Is” and “has” can be clearer than “serves as,” “stands as,” or “boasts,” while “can process” states the same relation more directly than “has the ability to process.” A more forceful verb helps only when it describes a more specific action.

3.2 Attribution determines checkability

Phrases such as “experts argue” and “studies show” borrow authority from unnamed sources. Once an expert, study, dataset, or observation is identified, the reader can inspect the basis of the claim. Without such support, the passage can state only the narrower fact established by the evidence.

Lists of media outlets create a related effect when they substitute notability for substance. A citation matters because of the claim it supports. One source tied to a specific claim is therefore more informative than a catalog of publication names.

3.3 Mechanisms replace ceremonial interpretation

Trailing participles such as “highlighting,” “showcasing,” and “underscoring” can append an interpretation without explaining why it follows. A mechanism, comparison, or consequence supplies the missing relation.

Agent-like draft

The interface uses blue and green, reflecting the team’s commitment to trust and growth while fostering a more engaging experience.

Reader-centered revision

The interface uses blue for selected controls and green for successful actions.

The revision assigns each color a function. If the source establishes only the colors, even that explanation would be too strong; the supported description would then end after “blue and green.”

3.4 Concrete endings preserve scope

Generic conclusions promise a bright future, celebrate a journey, or declare a step in the right direction. They change the mood without changing the reader’s understanding. A supported result often provides the natural ending, especially when its final clause identifies a limitation. In a document organized around a decision, closure may instead come from stating the decision that remains.

4 Local fluency can conceal relational errors

A phrase can sound natural while the sentence as a whole remains stiff or describes the wrong relation. Phrase substitution alone does not resolve this problem. The subject, verb, and modifiers have to form a coherent account of who acts, what changes, and under which conditions.

4.1 Question wrappers delay direct claims

A recurring agent-like construction introduces a known relation with a question word even though no question remains open. A noun or declarative clause can bring the answer into view earlier.

Agent-like draft

Each contribution should state what the study learned. The next section explains how the controller filters sensor noise.

Reader-centered revision

Each contribution should state the outcome of the study. The controller filters sensor noise by averaging five readings.

The first revision names the object directly, while the second places the filtering action in the main clause and its mechanism in a modifying phrase. A question-word clause remains appropriate when it expresses real uncertainty. “We do not know why the controller fails” cannot be replaced with a cause until the evidence identifies one.

4.2 Mechanical relative clauses distort purpose

Relative clauses can be grammatical and still force a strained relation. The strain appears in a sentence that makes the result seem to demand its own setup:

Agent-like draft

Include only the setup that is required by the result.

Reader-centered revision

Include only the setup needed to interpret the result.

The revision states the purpose of the setup without personifying the result. A reader-centered version, “Give readers only the setup they need to interpret the result,” shifts the emphasis again. The preferable construction depends on the intended focus, not on preserving active or passive voice as an end in itself.

4.3 Agency determines the useful subject

Active voice helps when responsibility or mechanism matters. “The system preserves the results automatically” identifies an actor that disappears from “The results are preserved automatically.” Likewise, “You do not need a configuration file” supplies the missing subject in “No configuration file needed.”

Passive voice remains useful when the procedure or affected object deserves attention, and it may be necessary when the actor is unknown. The choice between active and passive voice therefore depends on information structure. A sentence can keep the affected object in subject position while still naming the actor when responsibility affects the claim.

4.4 Packed catalogs hide governing relations

A comma series can conceal the relation among ideas even when it contains only three objects. Fixed taxonomies, procedures, and option sets are different because readers need to see every member. An ordinary inventory has no such justification.

Several inventory sentences in one paragraph flatten distinctions further. The syntax makes every item appear equally important, and a paragraph break does not change the underlying rhythm.

Agent-like draft

The report addresses hiring, onboarding, retention, scheduling, purchasing, and reporting.

Reader-centered revision

The report separates staffing from daily operations. Staffing covers hiring, onboarding, and retention, while daily operations include scheduling, purchasing, and reporting.

The revision names two categories, explaining why the six terms belong in the same passage and giving the writer stable labels for later reference. The parallel lists are informative because each instantiates one side of that division.

4.5 Stacked hedges blur calibrated uncertainty

“Could potentially possibly be argued” contains several weak signals but no clearer account of confidence. One well-chosen qualifier usually suffices. “May” marks possibility, while “likely” expresses greater confidence. “Under these conditions” restricts scope; “the available data suggest” identifies the evidential basis. Additional words do not produce a finer calibration unless they add another evidential distinction.

5 Visible organization can flatten priority

A document can look orderly while making its central claim harder to find. Agent-like answers can acquire headings and bullets before the dependencies among their ideas are clear. Once every item receives the same visual weight, the reader must reconstruct which claims support, qualify, or merely illustrate the others.

5.1 Lists clarify coherent sets

Lists work well when readers need to scan a fixed set or follow a sequence. Procedures and option sets often meet this condition. Connected claims usually belong in prose because their relationships matter more than their count.

A five-item threshold is useful only as a diagnostic; it has no mathematical status. Long lists can contain several unnamed groups, and short lists can still combine objects that serve unrelated functions. The form helps only when a governing relation makes the set coherent.

Forced triads create a subtler version of the problem. A sentence may group three claims merely because the cadence sounds complete, even though the subject has no three-part structure. False ranges have a similar effect. “From hiring to reporting” implies a scale absent from the underlying topics; “hiring and reporting” simply names them.

5.2 Headings reveal information architecture

A heading marks a real change in question, object, or level of analysis. One-line sections and sentences that repeat their headings create the appearance of navigation without reducing uncertainty. Bold label-and-colon bullets can make ordinary prose resemble a dashboard for the same reason.

Announcements such as “Let’s dive in” or “Here’s what you need to know” delay the content promised by the heading. Diff-anchored documentation similarly draws attention to the latest edit instead of describing the current system. Version-scoped documents are exceptions because change is their subject.

5.3 Cadence reflects the structure of the argument

Cadence depends on information order and clause structure as much as sentence length. Familiar context can appear near the opening, while an important result receives the stress of syntactic closure. A short sentence can emphasize a conclusion, but a run of clipped declarations gives every idea the same dramatic weight. Medium-length sentences can sound just as mechanical when each one ends with the same grammatical shape.

Agent-like draft

Then the new model arrived. No warning. No compromise. The old rules were gone.

Reader-centered revision

The new model changed the comparison because it did not share the assumptions of the earlier systems.

The revision joins the event with its explanation, so the causal relation determines the rhythm. A repeated pattern such as “treats X as Y, not Z” can be equally mechanical when it appears across neighboring paragraphs. Scientific distinctions often require contrast, especially when the contrast prevents an overclaim, but no single template fits all of them.

Slogan formulas manufacture a different kind of stress. “Symmetry is the language of trust” sounds complete while leaving the relation unspecified. A claim about symmetric layouts making controls easier to predict identifies an effect that evidence could support or refute.

5.4 Typography cannot establish provenance

Mechanical boldface, title case in every heading, decorative emoji, and inline-header lists can accompany agent-like prose. They become editing problems when they obscure hierarchy or conflict with the publication, not because any one of them reveals an author.

Punctuation also reflects the venue’s conventions and the writer’s established preferences. Humanizer removes dashes because its catalog treats them as a common default rhythm device; the broader general-writing standard permits an occasional dash when it matches the author’s sample and improves the sentence. Curly quotation marks are likewise a typographic convention. Neither mark has much forensic value by itself.

5.5 Stable names reduce memory demands

Synonym cycling can make one object appear to change. A paragraph that rotates among “the model,” “the system,” “the framework,” and “the approach” requires the reader to test whether each phrase has a new referent. Repeating the canonical term is quieter and more precise; variation becomes useful only when it expresses a real distinction.

6 Voice emerges from situated attention

I find situated attention the most useful distinction in this article because it directs analysis toward the choices that shape a passage. A writer notices the detail, tension, or exception that matters to this subject and this reader, then allocates space accordingly. Those judgments about relevance shape voice more deeply than ornament does.

6.1 Supported specifics resist generic paraphrase

Names, measurements, dates, odd quotations, and concrete observations narrow the range of plausible continuations. They also give the reader something to verify. Agent-like prose can round such a detail into broad significance, and an overzealous edit can erase the detail again even as it improves the flow.

Specificity does not license invention. An anecdote or sensory detail is grounded in the writer's experience or the available evidence. An opinion may express the writer's judgment without external proof, but an editor cannot invent that judgment. Quantitative claims remain bound to the evidence.

6.2 Unresolved tension can remain unresolved

People often hold mixed views. A writer may consider a policy broadly useful and still distrust one consequence; a study may show a strong result in one regime alongside an unexplained failure in another. When the evidence does not settle the tension, preserving it is more accurate than extracting a neat lesson.

An agent-generated continuation may turn uncertainty into a balanced conclusion, answer criticism with an optimistic final paragraph, or generalize from an unresolved observation. Leaving the tension intact is often more accurate because the evidence has not supplied the resolution that a balanced ending suggests.

6.3 Deliberate irregularity carries information

A distinctive word, abrupt transition, or parenthetical aside may look inefficient in isolation. Its value depends on whether the writer can account for the effect: perhaps it registers surprise, changes the level of confidence, or lets a technical caveat interrupt a broader claim. Deliberate irregularity carries voice when the reader can infer its purpose. Without such a purpose, the variation is simply random.

Era-bound references, local language, and field-specific phrasing can also mark voice. Their intelligibility depends on the audience. An explanation may be necessary for a broad readership, but replacement with generic internet language usually strips away the distinction they expressed.

6.4 Polish and provenance are different variables

Perfect grammar can come from a human author or editor. Formal vocabulary, a dry register, letter-style greetings, and common transitions are similarly ambiguous. One dash or one short sentence says little, and unsourced prose is a research problem rather than an authorship test. Templates account for much of the regularity in correctly formatted documents.

Secondhand language requires separate treatment because a suspicious phrase inside a quotation or title belongs to its source. Proper names and worked examples have the same protection. Their wording can be explained or cited, but revision would corrupt the evidence under discussion.

Version history supplies stronger evidence than stylistic guesswork. Text committed before ChatGPT's public release on November 30, 2022, could not have been produced with ChatGPT. Drafts and revision records provide similar provenance and stronger evidence than a detector's impression of the final surface.

6.5 Genre shapes the expression of personality

Essays and personal writing can support humor, first person, and openly mixed feelings. Technical or legal prose may require a neutral voice, yet neutrality need not become generic. Humanization works within those constraints by preserving the situated choices available in the genre instead of adding personality indiscriminately.

7 Possible causes of recurring agent-like patterns

Output patterns do not identify their own causes. I treat the explanations below as hypotheses about different stages of generation, not as settled accounts of why a model writes as it does. A deployed language agent reflects both pretraining and any subsequent post-training. At generation time, its behavior also depends on system instructions and the decoding procedure. Retrieval or tool use can change the evidence available to the model. Several of these mechanisms may operate together.

7.1 Local prediction may weaken larger structures

During pretraining, a language model learns to predict each next token from the available preceding context. Later training changes the model's behavior, but familiar phrases and locally smooth word sequences may remain attractive even when the sentence or paragraph requires a less common construction.

Local plausibility does not guarantee that a relation will remain clear across the full sentence. This pressure could produce stock transitions and repeated rhetorical frames; a penalty against lexical repetition could produce the opposite surface pattern, synonym cycling. The outputs look different, but in both cases a locally reasonable choice weakens the larger structure.

This account predicts that an explicit paragraph plan will reduce relation-level errors without changing the content to be expressed. A matched evaluation could compare outputs with and without that plan, measuring factual preservation and the clarity of the resulting relations. Such a result would test a behavioral consequence of the hypothesis, not identify the training stage that caused it.

7.2 An underspecified audience may produce generic prose

A prompt can specify a task without giving the model a stable representation of the person who will use the answer. In that case, the response may default to generic forms of helpfulness: it explains broadly, announces transitions, and labels ideas “important” without selecting the consequence that matters in context.

Readerless corporate shorthand fits this account. “Give each section one job” transfers across genres because it presupposes almost nothing. “Each section should help a reviewer answer one question” depends on a reader with a particular decision and may therefore produce a different selection of detail.

Testing this account requires more than tone labels. A matched comparison could hold subject and length constant while varying the reader's prior knowledge and pending decision. Members of the named audience could then evaluate whether each version supports the stated use.

7.3 Mixed training data may favor a middle register

Training corpora span genres, institutions, eras, and levels of editorial quality. Without a strong voice cue, a model may settle on language acceptable across much of that mixture. Common evaluative words and regular paragraph shapes are relatively portable, whereas unusual specifics depend more heavily on context.

Code and API documentation may exert a more particular influence, as may issue trackers and technical manuals. These genres rely on modular headings and explicit labels; they also make edge cases and procedures easy to scan. Transferred into an essay, the same conventions can yield tiny sections and prose that resembles an interface. This is a conjecture about cross-genre transfer, not a claim that technical documentation is poorly written.

Style conditioning offers one test. Output with no sample could be compared with output conditioned on one sample and then on several samples from the same genre. Differences among models with documented training emphases would provide another, less controlled comparison.

7.4 Post-training may favor answers that look complete

Instruction tuning and preference optimization may favor features that evaluators readily interpret as helpful. In a chat interface, headings, long lists, recaps, balanced caveats, and offers to continue can look thorough even when they weaken a finished document. The same incentives may reward automatic agreement: “Great question” and “You’re absolutely right” preserve rapport without testing the underlying claim.

Stacked hedges can also present the appearance of caution while leaving the actual uncertainty unspecified. The cleanest comparison would pair a base model with an instruction-tuned model from the same family on a matched continuation task. A second factor, chat reply versus finished document, would test whether the model can condition those habits on genre.

7.5 Pressure to answer may encourage unsupported claims

Some prompts demand an answer that the supplied evidence cannot support. Vague attribution and plausible gap-filling allow generation to continue, and the model may turn the missing evidence into a familiar story. A private person with little public documentation, for example, can acquire an invented preference for privacy.

A direct test would vary retrieval and whether the model may answer “unknown,” then check every factual claim against the evidence supplied in that condition. Fluency would not compensate for an unsupported claim.

These explanations can reinforce one another. A familiar phrase supplied by local prediction may survive because mixed training data make it broadly acceptable. Post-training can preserve it when it resembles a complete answer, and an underspecified reader gives the model no reason to remove irrelevant material. Experiments that vary one factor at a time would clarify which mechanisms affect each observed pattern.

8 Revision as constrained inference

Style revision is safest when it proceeds from the document’s evidence and reader model toward its sentences. The order matters because a local improvement can quietly change a factual claim, while a broad cleanup can erase the cadence that made the passage recognizably the writer’s.

8.1 The source sets the factual boundary

A conservative revision starts from the draft’s central point and intended reader. Its factual boundary and recurring voice constrain the possible rewrites; calibrated uncertainty does as well. Changing “suggests” to “shows” changes the claim even when the sentence becomes shorter.

8.2 Paragraphs expose the hierarchy of ideas

Clusters of formulaic language often reveal a paragraph whose governing relation is missing. At paragraph scale, the editor can identify the decisive claim and group evidence according to the role it plays. Evidence for a result, qualifications that specify its extent, and mechanistic details that explain why it may hold play different roles. Ordering them accordingly lets readers recover those dependencies without assigning every sentence the same visual or rhythmic weight.

8.3 Sentences expose agency and relation

Sentence-level analysis identifies the grammatical subject and verb, then tests the logical role of each modifier. At this scale, unnamed actors, mechanical relative clauses, and mismatched verbs become visible. Cadence also requires attention here, although the relevant pattern usually spans several sentences rather than one.

8.4 Comparison with the source reveals losses

Comparing the revised passage with the source may reveal a specific detail, qualification, or abrupt transition lost during cleanup. It also helps distinguish chat residue and decorative endings from genuine voice. Reading at speaking pace makes repeated clause shapes audible, but the source remains the best evidence for the writer's deliberate irregularities.

8.5 A compact example

Agent-like draft

Effective writing is not just about clarity; it is about fostering connection, enhancing engagement, and ensuring that ideas resonate. By leveraging a reader-centered approach, writers can create robust communication that truly makes an impact.

Reader-centered revision

Effective writing gives a particular reader the information needed to understand or use an idea.

The first draft combines a canned contrast with a forced group of three. Promotional language and a corporate verb add ceremony to a simple claim, while the reference to readers never identifies their need. The revision states that need directly.

This degree of compression would be wrong if the paragraph contained evidence, scope, or a mechanism. It works here because the discarded phrases repeat one unsupported abstraction. Revision can remove that repetition without establishing a general preference for one-sentence paragraphs.

The limiting case is an edit so regular that the writer disappears. Success is better measured by whether the writer's facts, judgment, and recognizable cadence remain, even where another editor would have made a different choice.

A Pattern taxonomy

This appendix preserves the complete Humanizer catalog in compact form. The numbers match version 2.15.0 of the skill. Each entry describes a recurrent surface pattern and the relation that a grounded revision recovers; the main text develops the principles shared across entries.

A.1 Inflated or unsupported content

1. **Undue emphasis on significance.** Claims about pivotal moments, legacies, and broader trends displace the event and its supported consequence.
2. **Notability and media lists.** Publication names function as borrowed importance when the passage never explains which cited claim is relevant.
3. **Superficial *-ing* analysis.** Trailing participles such as “highlighting” and “showcasing” append interpretation without a mechanism, comparison, or result.
4. **Promotional language.** Travel-copy adjectives, superlatives, and verbs such as “boasts” substitute evaluation for an observable property.
5. **Vague attribution.** “Experts argue” and “studies show” invoke an apparent source while leaving the claim impossible to check.
6. **Formulaic challenge and outlook sections.** A generic obstacle followed by optimism creates closure without identifying the specific constraint, response, or unresolved problem.

A.2 Lexical and grammatical defaults

7. **Overused agent vocabulary.** Recurrent words such as “pivotal,” “intricate,” “landscape,” “tapestry,” and “underscore” often carry tone after their descriptive content has disappeared.
8. **Copula avoidance.** “Serves as,” “stands as,” “features,” and “boasts” can make a simple relation less direct than “is” or “has.”
9. **Negative parallelism and tailing negation.** “Not only,” “not just,” and clipped endings such as “no guessing” turn a positive claim into a reusable contrast frame.
10. **Rule-of-three overuse.** A triad sounds complete even when the subject does not have three relevant parts.
11. **Synonym cycling.** Rotating near-synonyms for one stable object creates the appearance of distinctions that the argument does not make.
12. **False ranges.** “From X to Y” implies a scale or progression when X and Y are merely separate topics.
13. **Passive voice and subjectless fragments.** Agency disappears when a responsible actor matters, while fragments such as “No configuration needed” can leave the grammatical subject implicit.

A.3 Punctuation and visible formatting

14. **Dash-heavy rhythm.** Habitual dashes assign the same parenthetical stress across unrelated sentences. A writer’s established use of dashes remains relevant counterevidence.

15. **Mechanical boldface.** Repeated emphasis creates a visual hierarchy that may have no counterpart in the argument.
16. **Inline-header lists.** Bold label-and-colon bullets make connected prose resemble a set of independent interface fields.
17. **Title case everywhere.** Uniform title case makes every heading equally prominent even when the venue or hierarchy calls for a quieter convention.
18. **Decorative emoji.** Icons attached to headings or bullets add salience without carrying information.
19. **Curly quotation marks as a default tell.** Curly marks arise from typographic systems as readily as from language models and therefore carry almost no provenance information on their own.

A.4 Chat residue and missing evidence

20. **Collaborative chat artifacts.** “Of course,” “I hope this helps,” and offers to continue preserve dialogue inside a document that no longer has a conversational turn structure.
21. **Knowledge-cutoff disclaimers and gap-filling.** A missing fact becomes a paragraph about unavailable information, sometimes followed by a plausible story unsupported by the sources.
22. **Sycophantic tone.** Praise for the question or automatic agreement establishes rapport without evaluating the substance.

A.5 Filler, hedging, and endings

23. **Filler phrases.** “In order to,” “due to the fact that,” and “has the ability to” delay relations that “to,” “because,” and “can” express directly.
24. **Excessive hedging.** Several weak qualifiers obscure which single modal or condition matches the evidence.
25. **Generic positive conclusions.** Bright futures, journeys, and steps in the right direction add mood after the final supported result, limitation, or decision.
26. **Uniform hyphenation.** A compound remains hyphenated in predicate position even when ordinary grammar or the publication’s convention would drop the hyphen.

A.6 Rhetorical and structural formulas

27. **Persuasive authority tropes.** “The real question” and “at its core” claim interpretive authority before stating the question or relation.
28. **Signposting announcements.** “Let’s dive in” and “Now we will explore” announce content that could already occupy the sentence.
29. **Generic, fragmented, or overbuilt headings.** Vague display text, heading restatements, and tiny sections create visible structure without clarifying the argument.
30. **Diff-anchored writing.** Descriptions of the latest edit make a stable document depend on a version transition that matters only in changelogs, release notes, and migration guides.

31. **Manufactured punchlines and staccato drama.** Stacked fragments assign dramatic closure to each line rather than explaining the relation that makes a change important.
32. **Aphorism formulas.** Constructions such as “X is the language of Y” offer memorability while leaving the underlying effect unspecified.
33. **Conversational rhetorical openers.** “Honestly?” and similar fake-candid pauses delay an ordinary answer in order to stage a reveal.
34. **Question-word clause wrappers.** Non-question “what,” “how,” or “why” clauses postpone a known object or relation that could appear directly.
35. **Loose catalogs and list sprawl.** Series without a governing relation distribute equal weight across items whose roles differ. Adjacent inventory sentences preserve the same rhythm even across a paragraph break.

A.7 Sentence focus, reader purpose, and verb semantics

36. **Mechanical relative clauses.** A locally grammatical relative clause can force an unnatural passive or obscure the purpose of the head noun.
37. **Readerless corporate shorthand.** Metaphors of jobs, ownership, and deliverables, including the demand that a sentence “earn its place,” sound decisive while obscuring the passage’s actual relation to its reader. In practice, the passage may be supplying evidence for an inference or information needed for a decision.
38. **Semantically mismatched relations.** Claims that “stop” and sections that “know” receive actions their subjects cannot perform. Replacing the verb may still leave its object underspecified: conditions limit a claim’s extent, while a theorem’s consequences need not extend beyond a stated regime.
39. **Mechanical paragraph endings.** A paragraph fades out on setup or housekeeping, or appends a vague sentence merely to sound complete. Natural closure comes from the supported consequence, boundary, unresolved tension, or concrete fact at which the reasoning arrives.

A.8 Additional general-writing checks

General writing compares the broad cleanup with the source. This lets the editor restore a technical distinction, calibrated uncertainty, or recognizable cadence removed by an otherwise fluent revision. Specific facts and deliberate irregularities receive the same protection.

Binary contrasts, negative lists, throat-clearing, and faux-insight claims all delay a direct point, although a scientific contrast may remain essential to the claim. Colon reveals and rhetorical questions can manufacture drama, while repeated sentence shapes make unrelated ideas appear equivalent.

Fake-profound kickers and summary paragraphs create closure without advancing the argument. A quieter failure occurs when a paragraph ends on setup after its main point has already appeared. Anecdotes and detours behave differently when they establish context or tension. The distinction depends on the reading effect of the passage, not on a blanket preference for compression.

B Dimensions of comparative analysis

Several distinct dimensions organize the patterns in this article. Collapsing them into a single score would hide the difference between factual error, generic style, and a mismatch with one reader’s needs.

B.1 Semantic fidelity

The first dimension concerns preservation of meaning. A revision is faithful when its factual claims and examples remain grounded in the source, including statistics, quotations, citations, and technical distinctions. Calibrated uncertainty belongs to this dimension because a change in evidential verb can alter the claim even when every noun and number remains the same.

Verb semantics provide a related test. The subject must be able to perform the named action, literally or within an established technical convention, and its object or complement must fit the relation expressed by the verb. Evidence can support a claim; stated conditions may limit its extent or applicability. If the question concerns implication, the relevant boundary belongs to the consequences of the claim. Metaphorical movement or agency becomes a defect when it obscures one of these relations.

B.2 Reader model and information structure

The selection of detail in reader-centered prose reveals what the writer assumes about prior knowledge and intended use. At sentence scale, familiar context tends to appear early enough to connect with the preceding discourse, while the result or distinction receives an appropriate position of emphasis. At paragraph scale, evidence and qualifications remain visibly subordinate to the claim they support.

Several sentence patterns can disrupt this structure. Question-word wrappers delay the point, while distant subject-verb pairs postpone syntactic resolution. Unnamed actors and stiff relative clauses create different problems, so their importance depends on the reading effect rather than the presence of a particular token.

B.3 Rhetorical form

Generic significance, promotional language, vague attribution, decorative analysis, and optimistic conclusions share a common effect: they substitute rhetorical form for evidential support. Lists and headings cause a related problem when they distribute equal weight across ideas with different roles.

Display text belongs to the same analysis as body prose. Titles, abstracts, headings, captions, and example labels all frame an interpretation. A generic title or section-summary abstract can therefore misrepresent the argument even when every sentence in the body remains unchanged.

B.4 Cadence and voice

Cadence is visible in sentence length, but it also depends on openings, clause shapes, contrast formulas, and syntactic closure. Several short sentences may manufacture drama; several medium sentences with identical structures can be equally monotonous. Repeated formulas such as “treats X as Y, not Z” deserve attention when the syntax persists after the underlying comparison changes. Paragraph endings add another source of emphasis: a natural ending completes the current inference, whereas a generic summary or slogan only imitates closure.

Voice is not the opposite of regularity. It consists of supported specifics, recurring judgments, deliberate irregularities, and genre-appropriate personality. A successful revision reduces mechanical rhythm without normalizing these choices into the editor’s own style.

B.5 Provenance

None of the preceding dimensions establishes authorship. Quotations, proper names, and secondhand examples retain the provenance of their source, while templates and professional editing can explain a polished surface. Version history and drafting records provide stronger evidence about production

than stylistic inference. The patterns remain useful because they describe reading effects even when the author is known.

C Sources and acknowledgments

This article draws on five agent-facing resources in the [writing-skills](#) repository: Humanizer version 2.15.0, general-writing, better-usage, writing-cadence, and academic-voice.

Humanizer builds on Siqi Chen’s [humanizer](#), which draws its pattern catalog from Wikipedia’s [Signs of AI writing](#). The general-writing framework builds on Peter Yang’s [no-ai-slop](#). Both upstream projects are distributed under the MIT License, and their license notices remain with the corresponding skills. The cadence discussion also draws on [Purdue OWL’s account of sentence variety](#) and [Gopen and Swan’s analysis of reader expectations in scientific prose](#).

Better-usage supplies the analysis of semantically compatible subjects, verbs, and objects, including the need to name a missing intermediate relation. Academic-voice calls for prose that is careful about evidence, precise about relations, and addressed to a reader without adopting corporate or social-media shorthand.

The causal hypotheses in [Section 7](#) are conjectures presented for discussion and testing. These hypotheses are the author’s and do not represent positions taken by the upstream projects.